

Assignment 2

Kartikey Agarwal

Task Insights

Working through Assignment-2 offered hands-on experience with several core concepts in image processing and computer vision:

1 Task 1: Image Processing Fundamentals

- Captured real-time webcam data using OpenCV and applied various transformations in Python.
- Converted images between color spaces (RGB → HSV, LAB), which is useful for tasks like segmentation or analysis.
- Visualized individual channels such as Hue, Saturation, Value, and LAB components.
- Applied histogram equalization to enhance contrast in grayscale images, improving visual clarity.
- Explored edge detection using Laplacian and Scharr filters, observing how different kernels affect output.
- Used posterization and binary thresholding to simplify images for further processing.
- Added and reduced salt-and-pepper noise with a median filter, and performed sharpening using unsharp masking.

Key Python Functions

```
def hsv_split(img):  
    hsv = cv.cvtColor(img, cv.COLOR_BGR2HSV)  
    h, s, v = cv.split(hsv)  
    return h, s, v  
  
def histogram_equalization(img):  
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)  
    eq = cv.equalizeHist(gray)  
    return gray, eq
```

```

def binary_inversion(img):
    threshold_value = 127
    max_value = 255
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    ret, binary = cv.threshold(gray, threshold_value, max_value,
                               cv.THRESH_BINARY_INV)
    return binary

def posterize(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    posterized = (gray // 64) * 85
    return posterized

def edge_detection(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    blurred_image = cv.GaussianBlur(gray, (3, 3), 0)
    laplacian = cv.Laplacian(blurred_image, cv.CV_64F)
    laplacian_abs = cv.convertScaleAbs(laplacian)
    Gx = cv.Scharr(gray, cv.CV_64F, 1, 0)
    Gy = cv.Scharr(gray, cv.CV_64F, 0, 1)
    gradient_magnitude = cv.magnitude(Gx, Gy)
    return laplacian_abs, gradient_magnitude

def median_filter_denoise(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    noisy = gray.copy()
    h, w = gray.shape
    num_noise = int(h * w * 0.05)
    for _ in range(num_noise):
        i = random.randint(0, h - 1)
        j = random.randint(0, w - 1)
        noisy[i, j] = 255 if random.random() < 0.5 else 0
    denoised = cv.medianBlur(noisy, 3)
    return noisy, denoised

def unsharp_masking(img):
    amount = 1.5
    blur_ksize = (5, 5)
    sigma = 1.0
    blur = cv.GaussianBlur(img, blur_ksize, sigma)
    sharpened = cv.addWeighted(img, 1 + amount, blur, -amount, 0)
    return sharpened

def convert_to_lab(img):
    lab = cv.cvtColor(img, cv.COLOR_BGR2Lab)
    l, a, b = cv.split(lab)

```

```
return l, a, b
```

The screenshot shows a Jupyter Notebook interface with the following code in a cell:

```
import cv2 as cv
cap = cv.VideoCapture(0)

if not cap.isOpened():
    print("Error: Could not open webcam")
else:
    ret, frame = cap.read()
    if ret:
        flipframe = cv.flip(frame, 1)
        cv.imwrite("captured_img.jpg", flipframe)
        print("Image captured and saved as 'captured_img.jpg'")
    else:
        print("Error capturing image")
cv.destroyAllWindows()
cv.waitKey(0)
cap.release()
cv.destroyAllWindows()
```

The cell output shows a timestamp of 6.1s and the message "Image captured and saved as 'captured_img.jpg'". The notebook also lists other files in the workspace.

The screenshot shows a Jupyter Notebook interface with the following code in a cell:

```
img_cv.imread("captured_img.jpg")

def hsv_split(img):
    hsv_cv.cvtColor(img, cv.COLOR_BGR2HSV)
    h, s, v_cv.split(hsv)
    return h,s,v

def histogram_equalization(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    eq = cv.equalizeHist(gray)
    return gray, eq

def binary_inversion(img):
    threshold_value = 127
    max_value = 255
    gray_cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    ret, binary= cv.threshold(gray, threshold_value, max_value, cv.THRESH_BINARY_INV)
    return binary

def posterize(img):
    #dividing 256 levels in to 4 bins i.e. into 4 evenly spaced intensity levels
    #0-63: bin 0
    #64-127: bin 1
    #128-191: bin 3
    #192-255: bin 4
    # and then rescaling back to 0-255 range : (255/(4-1)*85) :: [0,85,170,255]
    gray_cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    posterized=(gray//64)*85
    return posterized

def edge_detection(img):
    gray_cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    blurred_image = cv.GaussianBlur(gray, (3, 3), 0)

    # Apply the Laplacian operator
    laplacian = cv.Laplacian(blurred_image, cv.CV_64F)
    # Convert the result to 8-bit (0-255) range
```

The cell output shows a timestamp of 0.0s. The notebook also lists other files in the workspace.

The screenshot shows the VS Code interface with the following details:

- File Explorer:** Shows a folder structure under "CODEFORCES" named "c".
- Code Editor:** Displays Python code for image processing tasks. The code includes functions for posterizing, edge detection, median filter denoising, unsharp masking, and converting to LAB color space.
- Terminal:** Shows command-line history related to image processing tasks.
- Status Bar:** Shows "Spaces: 4" and "Cell 5 of 5".

```
def posterize(img):
    # dividing 256 levels in to 4 bins i.e. into 4 evenly spaced intensity levels
    # 0-63: bin 0   64-127: bin 1   128-191: bin 2   192-255: bin 4
    # and then rescaling back to 0-255 range : (255/(4-1)*85) :: [0,85,170,255]
    gray=cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    posterized=(gray//64)*85
    return posterized

def edge_detection(img):
    gray=cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    blurred_image = cv.GaussianBlur(gray, (3, 3), 0)

    # Apply the Laplacian operator
    laplacian = cv.Laplacian(blurred_image, cv.CV_64F)
    # Convert the result to 8-bit (0-255) range

    laplacian_abs = cv.convertScaleAbs(laplacian)

    # Apply Scharr operator to find the x and y gradients
    Gx = cv.Scharr(gray, cv.CV_64F, 1, 0)
    Gy = cv.Scharr(gray, cv.CV_64F, 0, 1)

    # Compute the gradient magnitude
    gradient_magnitude = cv.magnitude(Gx, Gy)
    return laplacian_abs,gradient_magnitude

def median_filter_denoise(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    noisy = gray.copy()
    h, w = gray.shape
    num_noise = int(h * w * 0.05)

    for i in range(num_noise):
        i = random.randint(0, h - 1)
        j = random.randint(0, w - 1)
        noisy[i, j] = 255 if random.random() < 0.5 else 0

    denoised = cv.medianBlur(noisy, 3)
    return noisy, denoised

def unsharp_masking(img):
    amount=1.5
    blur_ksize=(5,5)
    sigma=1.0
    blur = cv.GaussianBlur(img, blur_ksize, sigma)
    sharpened = cv.addWeighted(img, 1-amount, blur, -amount, 0)
    return sharpened

def convert_to_lab(img):
    lab = cv.cvtColor(img, cv.COLOR_BGR2Lab)
    l, a, b = cv.split(lab)
    return l, a, b

#LAB is a three-dimensional, device-independent color space with axes for lightness (L*), green-red (a*), and blue-yellow (b*)
#It separates brightness from color, is widely used for accurate color measurement, and covers all colors visible to humans.
#The L* channel looks like a black-and-white image, containing all the brightness information
#The a* channel visualizes red vs. green; the b* channel visualizes yellow vs. blue.
#Color axis: L: brightness only, no colour      a: red-positive, green-negative      b: blue-negative, yellow-positive

def display_results(results):
    titles = [
        'Hue, Saturated and Value Channel',
        'Original Gray and Equalized',
        'Binary Inversion', 'Posterized (4-level)',
        'Laplacian / Scharr Edges',
        'Salt and Pepper Denoising', 'Sharpening', 'LAB color space'
    ]
```

The screenshot shows the VS Code interface with the following details:

- File Explorer:** Shows a folder structure under "CODEFORCES" named "c".
- Code Editor:** Displays Python code for image processing tasks. The code includes functions for posterizing, edge detection, median filter denoising, unsharp masking, and converting to LAB color space.
- Terminal:** Shows command-line history related to image processing tasks.
- Status Bar:** Shows "Spaces: 4" and "Cell 5 of 5".

```
def posterize(img):
    # dividing 256 levels in to 4 bins i.e. into 4 evenly spaced intensity levels
    # 0-63: bin 0   64-127: bin 1   128-191: bin 2   192-255: bin 4
    # and then rescaling back to 0-255 range : (255/(4-1)*85) :: [0,85,170,255]
    gray=cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    posterized=(gray//64)*85
    return posterized

def edge_detection(img):
    gray=cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    blurred_image = cv.GaussianBlur(gray, (3, 3), 0)

    # Apply the Laplacian operator
    laplacian = cv.Laplacian(blurred_image, cv.CV_64F)
    # Convert the result to 8-bit (0-255) range

    laplacian_abs = cv.convertScaleAbs(laplacian)

    # Apply Scharr operator to find the x and y gradients
    Gx = cv.Scharr(gray, cv.CV_64F, 1, 0)
    Gy = cv.Scharr(gray, cv.CV_64F, 0, 1)

    # Compute the gradient magnitude
    gradient_magnitude = cv.magnitude(Gx, Gy)
    return laplacian_abs,gradient_magnitude

def median_filter_denoise(img):
    gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
    noisy = gray.copy()
    h, w = gray.shape
    num_noise = int(h * w * 0.05)

    for i in range(num_noise):
        i = random.randint(0, h - 1)
        j = random.randint(0, w - 1)
        noisy[i, j] = 255 if random.random() < 0.5 else 0

    denoised = cv.medianBlur(noisy, 3)
    return noisy, denoised

def unsharp_masking(img):
    amount=1.5
    blur_ksize=(5,5)
    sigma=1.0
    blur = cv.GaussianBlur(img, blur_ksize, sigma)
    sharpened = cv.addWeighted(img, 1-amount, blur, -amount, 0)
    return sharpened

def convert_to_lab(img):
    lab = cv.cvtColor(img, cv.COLOR_BGR2Lab)
    l, a, b = cv.split(lab)
    return l, a, b

#LAB is a three-dimensional, device-independent color space with axes for lightness (L*), green-red (a*), and blue-yellow (b*)
#It separates brightness from color, is widely used for accurate color measurement, and covers all colors visible to humans.
#The L* channel looks like a black-and-white image, containing all the brightness information
#The a* channel visualizes red vs. green; the b* channel visualizes yellow vs. blue.
#Color axis: L: brightness only, no colour      a: red-positive, green-negative      b: blue-negative, yellow-positive

def display_results(results):
    titles = [
        'Hue, Saturated and Value Channel',
        'Original Gray and Equalized',
        'Binary Inversion', 'Posterized (4-level)',
        'Laplacian / Scharr Edges',
        'Salt and Pepper Denoising', 'Sharpening', 'LAB color space'
    ]
```

```

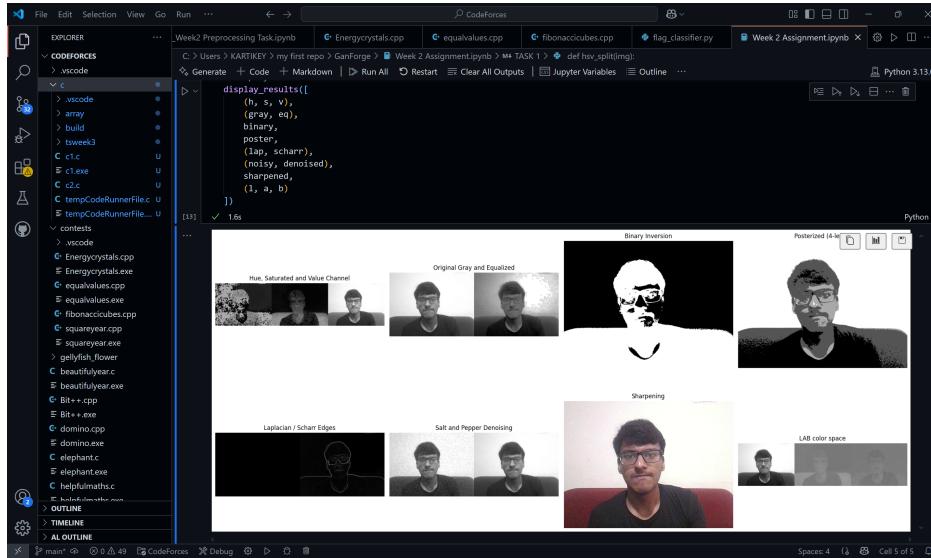
    fig, axes = plt.subplots(2, 4, figsize=(20, 10))
    for i, (title, images) in enumerate(zip(titles, results)):
        ax = axes[i//4, i%4] #rows=1//4 and columns=1%4
        if isinstance(images, tuple):
            # Combine side by side
            combined = np.hstack(images)
            cmap = "gray" if len(combined.shape) == 2 else None
        else:
            combined = images
            cmap = "gray" if len(combined.shape) == 2 else None
        ax.imshow(combined, cmap=cmap)
        ax.set_title(title)
        ax.axis('off')
    plt.tight_layout()
    plt.show()

# Convert image from BGR to RGB for consistent display
img_rgb = cv.cvtColor(img, cv.COLOR_BGR2RGB)

# Run all processing steps
h, s, v = hsv_split(img)
gray, eq = histogram_equalization(img)
binary = binary_inversion(img)
poster = posterize(img)
lap, schar = edge_detection(img)
noisy, denoised = median_filter_denoise(img)
sharpened = unsharp_masking(img_rgb)
i, a, b = convert_to_lab(img)

# Display results
display_results([
    (h, s, v),
    (gray, eq),
    (binary, poster),
    (lap, schar),
    (noisy, denoised),
    (i, a, b)
])

```



2 Task 2: YOLOv5 Object Detection

- Implemented object detection using the YOLOv5 model, following a tutorial.
- Tested the model on real-time webcam input and pre-recorded video.
- Detected various objects (persons, bottles, laptops, chairs, etc.).
- Model generally performed well, but sometimes misclassified or falsely detected objects (false positives).
- Some uncommon objects were not detected, likely due to limitations in the pre-trained dataset.

```
Anaconda Prompt
YOLOv5 summary: 4444 layers, 86705005 parameters, 0 gradients, 205.5 GFLOPs
video 1/1 (1/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 4 motorcycles, 1 bus, 1 truck, 1 traffic light, 1 bench, 2 backpacks, 1 chair, 1482.5ms
video 1/1 (2/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 4 motorcycles, 1 bus, 1 traffic light, 1 bench, 2 backpacks, 1 chair, 1472.0ms
video 1/1 (3/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 4 motorcycles, 1 bus, 2 traffic lights, 1 bench, 2 backpacks, 1 chair, 1440.5ms
video 1/1 (4/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 9 persons, 5 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 2 backpacks, 1399.0ms
video 1/1 (5/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 4 cars, 5 motorcycles, 2 buss, 1 truck, 2 traffic lights, 1 bench, 1 backpack, 1439.3ms
video 1/1 (6/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 5 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 2 backpacks, 1 chair, 1438.0ms
video 1/1 (7/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 7 persons, 1 bicycle, 4 cars, 4 motorcycles, 1 bus, 1 truck, 2 traffic lights, 1 bench, 2 backpacks, 1392.7ms
video 1/1 (8/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 7 persons, 1 bicycle, 4 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1392.4ms
video 1/1 (9/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 7 persons, 5 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1459.1ms
video 1/1 (10/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 7 persons, 5 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1357.9ms
video 1/1 (11/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 5 cars, 4 motorcycles, 1 bus, 3 traffic lights, 1 bench, 1 backpack, 1452.8ms
video 1/1 (12/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 4 motorcycles, 1 bus, 1 truck, 3 traffic lights, 1 bench, 1 backpack, 1445.8ms
video 1/1 (13/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 4 motorcycles, 1 bus, 3 traffic lights, 1 bench, 1 backpack, 1394.8ms
video 1/1 (15/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 5 cars, 6 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1411.0ms
video 1/1 (16/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 9 persons, 5 cars, 6 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1403.0ms
video 1/1 (17/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 10 persons, 5 cars, 7 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1404.3ms
video 1/1 (18/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 10 persons, 1 bicycle, 7 cars, 6 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1410.4ms
video 1/1 (19/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 1 bicycle, 7 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1399.0ms
video 1/1 (20/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 6 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1400.7ms
video 1/1 (21/301) C:\Users\KARTIKEY\Downloads\road_trafifc.mp4: 384x640 8 persons, 6 cars, 5 motorcycles, 1 bus, 2 traffic lights, 1 bench, 1 backpack, 1396.3ms
```





3 Task 3: Indonesia vs Poland Flag Classification

- Developed a Python program with OpenCV to classify an uploaded image as the flag of Indonesia or Poland.
- Both flags have two horizontal stripes (red and white), distinguished by the order:
 - **Indonesia:** Red on top, white on bottom
 - **Poland:** White on top, red on bottom
- Steps:
 1. Image preprocessing: Resize and crop to focus on the central portion.
 2. Color analysis: Convert to HSV, apply color masks to upper and lower halves.
 3. Color ratio comparison: Determine dominant color in each half to classify the flag.
- Classification worked well in clean conditions, but lighting, noise, and orientation could cause errors.
- A YOLO-based approach could improve robustness but would require a dedicated dataset.

The screenshot shows a Google Colab notebook titled "Week 2 Assignment.ipynb". The code in the cell is as follows:

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

def detect_and_crop_flag(image):
    """Detects red-white region and returns cropped flag image or None."""
    resized = cv.resize(image, (600, 400))
    hsv = cv.cvtColor(resized, cv.COLOR_BGR2HSV)

    # Red mask
    mask_red1 = cv.inRange(hsv, (0, 100, 100), (10, 255, 255)) # for Lower hue reds
    mask_red2 = cv.inRange(hsv, (160, 100, 100), (180, 255, 255)) # for Upper hue reds
    mask_red = cv.bitwise_or(mask_red1, mask_red2)

    # White mask
    mask_white = cv.inRange(hsv, (0, 0, 150), (180, 50, 255))
    combined_mask = cv.bitwise_or(mask_red, mask_white)

    # Find contours
    contours = cv.findContours(combined_mask, cv.RETR_EXTERNAL, cv.CHAIN_APPROX_SIMPLE)
    if not contours:
        return None, "No red-white region found."
    largest = max(contours, key=cv.contourArea)
    x, y, w, h = cv.boundingRect(largest)

    if w < 1.2 * h:
        return None, "Detected shape isn't flag-like."

    flag_crop = resized[y:y+h, x:x+w]
    return flag_crop, "flag detected and cropped."


def classify_flag(image):
    """Classifies flag image as Indonesia or Poland based on color in halves."""
    flag = cv.resize(image, (300, 200))
    hsv = cv.cvtColor(flag, cv.COLOR_BGR2HSV)

    # Top and bottom halves
    height = hsv.shape[0]
    top_half = hsv[:height//2, :]
    bottom_half = hsv[height//2:, :]

    # resize the image into width=100 and height=200 so top_half is taken for upper 100 pixels and bottom_half for lower 100 pixels
    # & mean = np.mean(..., axis=0) / value
    top_mean = np.mean(top_half, axis=(0, 1))
    bottom_mean = np.mean(bottom_half, axis=(0, 1))
    print(top_mean, bottom_mean)

    # Logic:
    # Indonesia: red (top), white (bottom)
    # Poland: white (top), red (bottom)

    # White in HSV: White color doesn't have hue, low Saturation (e.g., < 50) and high Value (e.g., > 150)
    # Red in HSV: Hue around 0-10 or 160-180, high Saturation (e.g., > 100), high Value

    # Condition for Indonesia: Red on top AND White on bottom
    # taking a safe range
```

The screenshot shows the continuation of the code from the previous snippet. The code in the cell is as follows:

```
if w < 1.2 * h:
    return None, "Detected shape isn't flag-like."

flag_crop = resized[y:y+h, x:x+w]
return flag_crop, "flag detected and cropped."


def classify_flag(image):
    """Classifies flag image as Indonesia or Poland based on color in halves."""
    flag = cv.resize(image, (300, 200))
    hsv = cv.cvtColor(flag, cv.COLOR_BGR2HSV)

    # Top and bottom halves
    height = hsv.shape[0]
    top_half = hsv[:height//2, :]
    bottom_half = hsv[height//2:, :]

    # resize the image into width=100 and height=200 so top_half is taken for upper 100 pixels and bottom_half for lower 100 pixels
    # & mean = np.mean(..., axis=0) / value
    top_mean = np.mean(top_half, axis=(0, 1))
    bottom_mean = np.mean(bottom_half, axis=(0, 1))
    print(top_mean, bottom_mean)

    # Logic:
    # Indonesia: red (top), white (bottom)
    # Poland: white (top), red (bottom)

    # White in HSV: White color doesn't have hue, low Saturation (e.g., < 50) and high Value (e.g., > 150)
    # Red in HSV: Hue around 0-10 or 160-180, high Saturation (e.g., > 100), high Value

    # Condition for Indonesia: Red on top AND White on bottom
    # taking a safe range
```

```

# Condition for Indonesia: Red on top AND white on bottom
# taking a safe range
is_top_red = (top_mean[0] < 20 or top_mean[0] > 150) and top_mean[1] > 100
is_bottom_white = bottom_mean[1] < 65 and bottom_mean[2] > 150 # Check saturation and value for white

if is_top_red and is_bottom_white:
    return "Indonesia Flag"

# Condition for Poland: White on top AND Red on bottom
# taking a safe range
is_top_white = top_mean[1] < 65 and top_mean[2] > 150
is_bottom_red = (bottom_mean[0] < 20 or bottom_mean[0] > 150) and bottom_mean[1] > 100

if is_top_white and is_bottom_red:
    return "Poland Flag"

else:
    return "Uncertain"

image=cv.imread("poland.jpg")
Flag_crop, msg = detect_and_crop_flag(image)
print(msg)
Flag_crop=cv.cvtColor(Flag_crop, cv.COLOR_BGR2RGB)
plt.imshow(Flag_crop_rgb)
plt.axis('off')
result = classify_flag(Flag_crop)
print("Classification: " + result)

```

Flag detected and cropped.
[106.8889333 67.3529333 244.71683333] [154.8932 150.92536667 198.49713333]
Classification: Poland Flag

```

return "Uncertain"

image=cv.imread("poland.jpg")
Flag_crop, msg = detect_and_crop_flag(image)
print(msg)
Flag_crop=cv.cvtColor(Flag_crop, cv.COLOR_BGR2RGB)
plt.imshow(Flag_crop_rgb)
plt.axis('off')
result = classify_flag(Flag_crop)
print("Classification: " + result)

```

Flag detected and cropped.
[106.8889333 67.3529333 244.71683333] [154.8932 150.92536667 198.49713333]
Classification: Poland Flag

4 Task 4: Understanding YOLOv5

Introduction

YOLOv5 is a popular model for object detection in images and videos. It processes the entire image in one pass, making it faster and suitable for real-time applications.

Key Components

- **Backbone (CSPDarknet):** Extracts features from the input image using Cross-Stage Partial connections.

- **Neck (PANet):** Aggregates features from different layers, enhancing detection at various scales.
- **Head:** Produces the final predictions (object locations and classes).

Training Techniques

- **Data Augmentation:** Mosaic augmentation, combining multiple images for better generalization.
- **Loss Functions:** Advanced loss functions balance localization, classification, and objectness.
- **Anchor Generation:** Uses anchor boxes to predict object sizes and shapes.
- **FP16 Precision:** Reduces memory and speeds up computation.

Transition to PyTorch

YOLOv5 moved from Darknet to PyTorch, making it more accessible and easier to integrate.

Performance and Variants

YOLOv5 comes in several sizes (n, s, m, l, x), offering trade-offs between speed and accuracy for different applications.

Conclusion

YOLOv5 is notable for its speed, accuracy, and ease of use, making it a preferred choice for real-time object detection tasks.

Github link: <https://github.com/Kari3020/GanForge>